

# Figure Skating Video Classification

by Elena Golubeva and Vladislav Goteiner



**01**

**ABOUT THE PROJECT**

**02**

**PROCESS**

**03**

**MODEL STRUCTURE**

**04**

**CODE EXAMPLE**


**05**

**FINAL RESULTS**

# FIGURE SKATING JUMP CLASSIFICATION

## About the project:


Tell what kind of jump is on the given video. What score would figure skater get for his element execution.



# FIGURE SKATING JUMP CLASSIFICATION

## Expectations:


To classify element with over 85% accuracy. Give a score based on similar element executions that we learned.



# FIGURE SKATING JUMP CLASSIFICATION

## Obstacles:

Examples from different classes has very little differences.




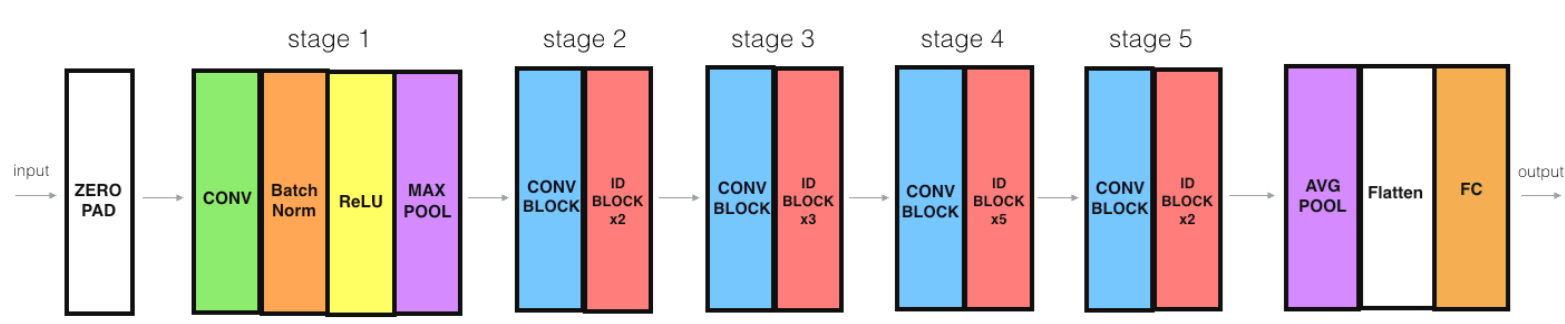
**CREATE MODEL**



**PREPARE TEST DATA**

**TRAIN ON THE MODEL**



**RESULT DISPLAY**












- ReLU (rectified linear unit)



$$f_{\text{ReLU}}(x) = \max(0, x)$$


ReLU can be approximated by **softplus** function

$$f_{\text{softplus}}(x) = \log(1 + e^x)$$




- The only non-linearity comes from the path selection with individual neurons being active or not
- It allows **sparse representations**:
  - for a given input only a subset of neurons are active








(a) Standard Neural Net



(b) After applying dropout.





```

parent = os.listdir("/content/drive/My Drive/lena_space/train_video")

for video_class in parent[0:]:
    print(video_class)
    listing = os.listdir("/content/drive/My Drive/lena_space/train_video/"
                        + video_class)
    count = 1
    for file in listing:
        video = cv2.VideoCapture("/content/drive/My Drive/lena_space/train_video/"
                                + video_class + "/" + file)

        frameId = video.get(1)
        framerate = video.get(5)
        os.makedirs("/content/drive/My Drive/lena_space/train_frames/"
                    + video_class + "/" + "video_" + str(int(count)))
        while video.isOpened():
            success, image = video.read()
            if success != True:
                break
            frame_count = 1
            while success:
                image = cv2.resize(image, (224, 224), interpolation=cv2.INTER_AREA)
                cv2.imwrite("/content/drive/My Drive/lena_space/train_frames/"
                            + video_class + "/" + "video_" + str(
                                int(count)) + "/image_%d.jpg" % frame_count, image)
                success, image = video.read()
                frame_count += 1
        video.release()
    count += 1

```

```

parent = os.listdir("/content/drive/My Drive/lena_space/train_frames")
data = []
labels = []
scores = []
tags=["Axel","Euler","Flip","Loop","Lutz","Salchow","Toeloop"]
tag=0
count = 0
for video_class in parent[0:]: # it also contains DS.store file
    child = os.listdir("/content/drive/My Drive/lena_space/train_frames/" + video_class)
    for class_i in child[0:]:
        sub_child = os.listdir("/content/drive/My Drive/lena_space/train_frames/"
                                + video_class + "/" + class_i)
        for image_fol in sub_child[0:]:
            if count % 2 == 0: # (selected images at gap of 4)
                image = cv2.imread("/content/drive/My Drive/lena_space/train_frames/"
                                    + video_class + "/" + class_i + "/" + image_fol)
                image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
                data.append(image)
                labels.append(tags[tag])
                scores.append(randrange(3, 6))
                count += 1
            tag += 1

```

```

labels = np.array(labels)
data = np.array(data)
scores = np.array(scores)
lb = LabelBinarizer()
labels = lb.fit_transform(labels)
(trainX, testX, trainY, testY) = train_test_split(data, labels, test_size=0.25,
                                                    stratify=labels, random_state=42)

# initialize the training data augmentation object
trainAug = ImageDataGenerator(
    rotation_range=30,
    zoom_range=0.15,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.15,
    horizontal_flip=True,
    fill_mode="nearest")

# initialize the validation/testing data augmentation object
#(which we'll be adding mean subtraction to)
valAug = ImageDataGenerator()

# define the ImageNet mean subtraction (in RGB order) and set the mean subtraction value
#for each of the data augmentation objects
mean = np.array([123.68, 116.779, 103.939], dtype="float32")
trainAug.mean = mean
valAug.mean = mean

```



```

# load the ResNet-50 network, ensuring the head FC layer sets are left off
baseModel = ResNet50(weights="imagenet", include_top=False, input_tensor=Input(shape=(224, 224, 3)))
# construct the head of the model that will be placed on top of the the base model
headModel = baseModel.output
headModel = AveragePooling2D(pool_size=(7, 7))(headModel)
headModel = Flatten(name="flatten")(headModel)
headModel = Dense(512, activation="relu")(headModel)
headModel = Dense(512, activation="relu")(headModel)
headModel = Dropout(0.5)(headModel)
headModel = Dense(len(lbl.classes_), activation="softmax")(headModel)
model = Model(inputs=baseModel.input, outputs=headModel)
for layer in baseModel.layers:
    layer.trainable = False
# compile our model (this needs to be done after our setting our layers to being non-trainable)
print("[INFO] compiling model...")
opt = SGD(lr=0.01, momentum=0.9, decay=0.001)
model.compile(loss="categorical_crossentropy", optimizer=opt, metrics=["accuracy"])
# train the head of the network for a few epochs (all other layers are frozen) --
# this will allow the new FC layers to start to become
# initialized with actual "learned" values versus pure random
print("[INFO] training head...")
H = model.fit_generator(
    trainAug.flow(trainX, trainY, batch_size=32),
    steps_per_epoch=len(trainX) // 32,
    validation_data=valAug.flow(testX, testY),
    validation_steps=len(testX) // 32,
    epochs=10)

```



```
Epoch 1/10
59/59 [=====] - 538s 9s/step - loss: 1.7998 - accuracy: 0.2959 - val_loss: 1.6521 - val_accuracy: 0.3520
Epoch 2/10
59/59 [=====] - 527s 9s/step - loss: 1.4022 - accuracy: 0.4758 - val_loss: 1.6664 - val_accuracy: 0.3680
Epoch 3/10
59/59 [=====] - 537s 9s/step - loss: 1.2163 - accuracy: 0.5599 - val_loss: 1.8830 - val_accuracy: 0.4191
Epoch 4/10
59/59 [=====] - 539s 9s/step - loss: 1.1058 - accuracy: 0.6078 - val_loss: 1.4093 - val_accuracy: 0.4043
Epoch 5/10
59/59 [=====] - 543s 9s/step - loss: 0.9647 - accuracy: 0.6594 - val_loss: 1.1594 - val_accuracy: 0.4290
Epoch 6/10
59/59 [=====] - 542s 9s/step - loss: 0.8872 - accuracy: 0.6770 - val_loss: 2.1557 - val_accuracy: 0.4010
Epoch 7/10
59/59 [=====] - 542s 9s/step - loss: 0.7931 - accuracy: 0.7172 - val_loss: 1.8893 - val_accuracy: 0.4521
Epoch 8/10
59/59 [=====] - 538s 9s/step - loss: 0.7264 - accuracy: 0.7385 - val_loss: 1.4750 - val_accuracy: 0.4620
Epoch 9/10
59/59 [=====] - 539s 9s/step - loss: 0.6473 - accuracy: 0.7504 - val_loss: 1.4693 - val_accuracy: 0.4505
Epoch 10/10
59/59 [=====] - 541s 9s/step - loss: 0.6132 - accuracy: 0.7845 - val_loss: 1.6687 - val_accuracy: 0.3944
```



```

print("[INFO] loading model and label binarizer...")
model = load_model('/content/drive/My Drive/lena_space/classification.model')
lb = pickle.loads(open('/content/drive/My Drive/lena_space/lb.pickle', "rb").read())
mean = np.array([123.68, 116.779, 103.939][::1], dtype="float32")
Q = deque(maxlen=10)

vs = cv2.VideoCapture('/content/drive/My Drive/lena_space/test_video/salchow.mp4')
writer = None
(W, H) = (None, None)
# loop over frames from the video file stream
while True:
    # read the next frame from the file
    (grabbed, frame) = vs.read()
    # if the frame was not grabbed, then we have reached the end of the stream
    if not grabbed:
        break
    # if the frame dimensions are empty, grab them
    if W is None or H is None:
        (H, W) = frame.shape[:2]
        # clone the output frame, then convert it from BGR to RGB ordering, resize the frame to a
        # fixed 224x224, and then perform mean subtraction
        output = frame.copy()
        frame = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
        frame = cv2.resize(frame, (224, 224)).astype("float32")
        frame -= mean
    # make predictions on the frame and then update the predictions queue
    preds = model.predict(np.expand_dims(frame, axis=0))[0]
    Q.append(preds)

```

```

# perform prediction averaging over the current history of previous predictions
results = np.array(Q).mean(axis=0)
i = np.argmax(results)
label = lb.classes_[i]
# draw the activity on the output frame
text = "activity: {}".format(label)
text_position = ((int) (image.shape[1]/2 - 268/2), (int) (image.shape[0]/2 - 36/2))
cv2.putText(output, text, text_position, (35, 50), cv2.FONT_HERSHEY_SIMPLEX,
            1.25, (112,128,144), 5)
# check if the video writer is None
if writer is None:
    # initialize our video writer
    fourcc = cv2.VideoWriter_fourcc(*"DIVX")
    writer = cv2.VideoWriter('/content/drive/My Drive/result_of_video_classification.avi',
                            fourcc, 30, (W, H), True)
# write the output frame to disk
writer.write(output)
# show the output image
cv2.imshow(output)
key = cv2.waitKey(1) & 0xFF
# if the `q` key was pressed, break from the loop
if key == ord("q"):
    break
# release the file pointers
print("[INFO] cleaning up...")
writer.release()
vs.release()

```

activity: Axel

3 Loop			
BV	4.90	GOE	0.14
C	12.08		

