

# Introduction to Reinforcement Learning Course - Final project

Dr. Teddy Lazebnik

1.10.2024

## Abstract

The final project aims to apply reinforcement learning techniques to a real-world or simulated problem. Students will independently conceive, design, implement, and evaluate a reinforcement learning solution. The project should demonstrate a deep understanding of the course material and the ability to apply it creatively.

## Project Requirements

### Project Selection

Students are expected to propose a unique and challenging reinforcement learning project that aligns with the course content and demonstrates their ability to apply the learned concepts.

### Project Topic Selection

**Relevance:** The project should be relevant to the field of reinforcement learning and have potential real-world applications. **Originality:** Students are encouraged to propose original project ideas or explore novel approaches to existing problems. **Feasibility:** The project should be achievable within the given timeframe and resources.

### Problem Formulation

**Clear Definition:** The reinforcement learning problem should be clearly defined, including: States: The set of all possible states the agent can be in; Actions: The set of actions the agent can take in each state; Rewards: The feedback signal that the agent receives for its actions; and Environment Dynamics: The rules governing the transitions between states. **Objective:** The project should have a well-defined goal or objective for the agent to achieve.

### Data and Environment

**Data Availability:** If the project involves real-world data, students should ensure its availability and accessibility. **Environment Creation:** If a simulated environment is required, students should consider using existing platforms or building their own. **Environment Complexity:** The chosen environment should provide sufficient challenges for the reinforcement learning agent.

### Project Scope

The project scope should demonstrate a comprehensive understanding of reinforcement learning concepts and their application to a real-world or simulated problem.

## Depth of Exploration

**Algorithm Selection:** Students should justify their choice of reinforcement learning algorithm(s) based on the problem characteristics and state-of-the-art research. **Hyperparameter Tuning:** A thorough exploration of hyperparameter space is expected to optimize agent performance. **Evaluation Metrics:** Students should define appropriate metrics to assess the agent's performance and compare different approaches.

## Innovation and Creativity

**Novel Approaches:** Students are encouraged to explore innovative techniques or combinations of methods. **Theoretical Contributions:** If applicable, students can contribute to the theoretical understanding of reinforcement learning. **Practical Impact:** The project should demonstrate the potential practical impact of the developed solution.

## Project Deliverables

The final project deliverables should comprehensively document the project process, results, and code.

### Written Report

- **Project Overview:** Clear and concise description of the project goals, motivation, and scope.
- **Problem Formulation:** Detailed explanation of the reinforcement learning problem, including state space, action space, reward function, and environment dynamics.
- **Methodology:** Description of the chosen reinforcement learning algorithm(s), implementation details, and hyperparameter tuning process.
- **Results:** Presentation of experimental results, including quantitative and qualitative analysis.
- **Discussion:** Analysis of the results, comparison to existing work, and identification of limitations.
- **Conclusion:** Summary of the project's findings and contributions.

### Code

- **Clean and Organized Code:** Well-structured, commented, and readable code implementation of the reinforcement learning agent.
- **Reproducibility:** Clear instructions on how to run the code and reproduce the results.
- **Version Control:** Use of version control (e.g., Git) to manage code changes.

### Presentation

- **Clear Communication:** Effective presentation of the project's key findings and insights - 15 minutes.
- **Visual Aids:** Use of slides, demos, or other visual aids to enhance understanding - your call.

## Evaluation Criteria

The final project will be evaluated based on the following criteria:

- **Problem Formulation:** The clarity and depth of the problem definition, including state space, action space, and reward function.

- Algorithm Selection: The appropriateness of the chosen reinforcement learning algorithm(s) for the problem.
- Implementation Quality: The correctness, efficiency, and maintainability of the code implementation.

### **Additional Comments**

Students must adhere to academic integrity guidelines. Plagiarism and collusion will not be tolerated. If working in groups, all members must contribute to the project. Groups of up to 4 students are allowed. Please note that larger groups are expected to produce “larger” projects and reports. The reports are expected to be between 8 and 20 pages long (not including the reference list) with 1.5 line spacing, Arial 12 font, one-inch margin.